

# Data Analysis Railway in UK

## Part 1

### Toolkit & Loading Data and Inspecting

#### Toolkit

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

#### Loading Data and Inspecting

```
df = pd.read_csv('/content/railway.csv')
```

```
df.head()
```

↗

	Transaction ID	Date of Purchase	Time of Purchase	Purchase Type	Payment Method	Railcard	Ticket Class	Ticket Type	Price	Departure Station	Arrival Destination	Date of Journey	Departure Time
0	da8a6ba8-b3dc-4677-b176	2023-12-08	12:41:11	Online	Contactless	Adult	Standard	Advance	43	London Paddington	Liverpool Lime Street	2024-01-01	11:00
1	b0cdd1b0-f214-4197-be53	2023-12-16	11:23:01	Station	Credit Card	Adult	Standard	Advance	23	London Kings Cross	York	2024-01-01	09:45
2	f3ba7a96-f713-40d9-9629	2023-12-19	19:51:27	Online	Credit Card	NaN	Standard	Advance	3	Liverpool Lime Street	Manchester Piccadilly	2024-01-02	18:15
3	b2471f11-4fe7-4c87-8ab4	2023-12-20	23:00:36	Station	Credit Card	NaN	Standard	Advance	13	London Paddington	Reading	2024-01-01	21:30
4	2be00b45-0762-485e-a7a3	2023-12-27	18:22:56	Online	Contactless	NaN	Standard	Advance	76	Liverpool Lime Street	London Euston	2024-01-01	16:45

```
df.info()
```

```
↗ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 31653 entries, 0 to 31652
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Transaction ID         31653 non-null object
1   Date of Purchase       31653 non-null object
2   Time of Purchase       31653 non-null object
3   Purchase Type          31653 non-null object
4   Payment Method         31653 non-null object
5   Railcard               10735 non-null object
6   Ticket Class           31653 non-null object
7   Ticket Type            31653 non-null object
8   Price                  31653 non-null int64
9   Departure Station      31653 non-null object
10  Arrival Destination    31653 non-null object
11  Date of Journey        31653 non-null object
12  Departure Time         31653 non-null object
13  Arrival Time           31653 non-null object
14  Actual Arrival Time    29773 non-null object
15  Journey Status         31653 non-null object
16  Reason for Delay       4172 non-null  object
17  Refund Request         31653 non-null object
dtypes: int64(1), object(17)
memory usage: 4.3+ MB
```

```
df.describe().round(4)
```



## Price

<b>count</b>	31653.0000
<b>mean</b>	23.4392
<b>std</b>	29.9976
<b>min</b>	1.0000
<b>25%</b>	5.0000
<b>50%</b>	11.0000
<b>75%</b>	35.0000
<b>max</b>	267.0000

pip install ydata-profiling



```
Requirement already satisfied: ydata-profiling in /usr/local/lib/python3.11/dist-packages (4.16.1)
Requirement already satisfied: scipy<1.16,>=1.4.1 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (1.15.3)
Requirement already satisfied: pandas!=1.4.0,<3.0,>1.1 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (2.2.2)
Requirement already satisfied: matplotlib<3.10,>=3.5 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (3.10.0)
Requirement already satisfied: pydantic>=2 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (2.11.7)
Requirement already satisfied: PyYAML<6.1,>=5.0.0 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (6.0.2)
Requirement already satisfied: Jinja2<3.2,>=2.11.1 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (3.1.6)
Requirement already satisfied: visions<0.8.2,>=0.7.5 in /usr/local/lib/python3.11/dist-packages (from visions[type_image_path]<0.8.2)
Requirement already satisfied: numpy<2.2,>=1.16.0 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (2.0.2)
Requirement already satisfied: htmlmin==0.1.12 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (0.1.12)
Requirement already satisfied: phik<0.13,>=0.11.1 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (0.12.5)
Requirement already satisfied: requests<3,>=2.24.0 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (2.32.3)
Requirement already satisfied: tqdm<5,>=4.48.2 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (4.67.1)
Requirement already satisfied: seaborn<0.14,>=0.10.1 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (0.13.2)
Requirement already satisfied: multimethod<2,>=1.4 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (1.12)
Requirement already satisfied: statsmodels<1,>=0.13.2 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (0.14.5)
Requirement already satisfied: typeguard<5,>=3 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (4.4.4)
Requirement already satisfied: imagehash==4.3.1 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (4.3.1)
Requirement already satisfied: wordcloud>=1.9.3 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (1.9.4)
Requirement already satisfied: dacite>=1.8 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (1.9.2)
Requirement already satisfied: numba<=0.61,>=0.56.0 in /usr/local/lib/python3.11/dist-packages (from ydata-profiling) (0.60.0)
Requirement already satisfied: PyWavelets in /usr/local/lib/python3.11/dist-packages (from imagehash==4.3.1->ydata-profiling) (1.9.6)
Requirement already satisfied: pillow in /usr/local/lib/python3.11/dist-packages (from imagehash==4.3.1->ydata-profiling) (11.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from Jinja2<3.2,>=2.11.1->ydata-profiling) (3.0.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (1.2.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (4.53.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (24.1)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (3.2.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib<3.10,>=3.5->ydata-profiling) (2.9.0)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in /usr/local/lib/python3.11/dist-packages (from numba<=0.61,>=0.56.0->ydata-profiling) (0.44.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas!=1.4.0,<3.0,>1.1->ydata-profiling) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas!=1.4.0,<3.0,>1.1->ydata-profiling) (2024.1)
Requirement already satisfied: joblib>=0.14.1 in /usr/local/lib/python3.11/dist-packages (from phik<0.13,>=0.11.1->ydata-profiling) (1.4.2)
Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.11/dist-packages (from pydantic>=2->ydata-profiling) (0.6.0)
Requirement already satisfied: pydantic-core==2.33.2 in /usr/local/lib/python3.11/dist-packages (from pydantic>=2->ydata-profiling) (2.33.2)
Requirement already satisfied: typing-extensions>=4.12.2 in /usr/local/lib/python3.11/dist-packages (from pydantic>=2->ydata-profiling) (4.12.2)
Requirement already satisfied: typing-inspection>=0.4.0 in /usr/local/lib/python3.11/dist-packages (from pydantic>=2->ydata-profiling) (0.4.0)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.24.0->ydata-profiling) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.24.0->ydata-profiling) (3.10.2)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.24.0->ydata-profiling) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests<3,>=2.24.0->ydata-profiling) (2024.12.14)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.11/dist-packages (from statsmodels<1,>=0.13.2->ydata-profiling) (0.5.6)
Requirement already satisfied: attrs>=19.3.0 in /usr/local/lib/python3.11/dist-packages (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2)
Requirement already satisfied: networkx>=2.4 in /usr/local/lib/python3.11/dist-packages (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2)
Requirement already satisfied: puremagic in /usr/local/lib/python3.11/dist-packages (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->matplotlib<3.10,>=3.5) (1.17.0)
```

```
from ydata_profiling import ProfileReport
ProfileReport(df)
```



Summarize dataset: 100%

29/29 [00:16<00:00, 2.67it/s, Completed]

0%| | 0/18 [00:00<?, ?it/s]  
6%| | 1/18 [00:01<00:18, 1.08s/it]  
100%| | 18/18 [00:12<00:00, 1.43it/s]

Generate report structure: 100%

1/1 [00:13<00:00, 13.59s/it]

Render HTML: 100%

1/1 [00:01<00:00, 1.73s/it]

## YData Profiling Report

Overview Variables Interactions Correlations Missing values Sample

# Overview

Brought to you by [YData](#)

Overview Alerts 10 Reproduction

Dataset statistics

Number of variables	18
Number of observations	31653
Missing cells	50279
Missing cells (%)	8.8%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	4.3 MiB
Average record size in memory	144.0 B

Variable types

Text	1
DateTime	6
Categorical	9
Numeric	1
Boolean	1

# Variables

Select Columns

## Part 2

## Data Cleaning and Preparation

### Handle missing values

```
df.isnull().sum()
```

	0
Transaction ID	0
Date of Purchase	0
Time of Purchase	0
Purchase Type	0
Payment Method	0
Railcard	20918
Ticket Class	0
Ticket Type	0
Price	0
Departure Station	0
Arrival Destination	0
Date of Journey	0
Departure Time	0
Arrival Time	0
Actual Arrival Time	1880
Journey Status	0
Reason for Delay	27481
Refund Request	0

dtype: int64

```
total_rows = len(df)
```

```
missing_railcard = df['Railcard'].isnull().sum()
```

```
missing_actual_arrival_time = df['Actual Arrival Time'].isnull().sum()
```

```
missing_reason_for_delay = df['Reason for Delay'].isnull().sum()
```

```
print(f"Missing values in 'Railcard': {missing_railcard} ({missing_railcard / total_rows:.2%})")
```

```
print(f"Missing values in 'Actual Arrival Time': {missing_actual_arrival_time} ({missing_actual_arrival_time / total_rows:.2%})")
```

```
print(f"Missing values in 'Reason for Delay': {missing_reason_for_delay} ({missing_reason_for_delay / total_rows:.2%})")
```

```
Missing values in 'Railcard': 20918 (66.09%)
Missing values in 'Actual Arrival Time': 1880 (5.94%)
Missing values in 'Reason for Delay': 27481 (86.82%)
```

```
df['Railcard'] = df['Railcard'].fillna('Unknown')
```

```
df['Reason for Delay'] = df['Reason for Delay'].fillna('Unknown Reason')
```

```
on_time_mask = df['Journey Status'] == 'On Time'
```

```
df.loc[on_time_mask, 'Actual Arrival Time'] = df.loc[on_time_mask, 'Actual Arrival Time'].fillna(df.loc[on_time_mask, 'Arrival Time'])
```

```
# Remove trailing space from 'Adult ' in 'Railcard' column
```

```
df['Railcard'] = df['Railcard'].replace('Adult ', 'Adult')
```

```
print("Percentage of missing values after handling:")
```

```
print(df.isnull().mean() * 100)
```

```
Percentage of missing values after handling:
```

Transaction ID	0.000000
Date of Purchase	0.000000
Time of Purchase	0.000000
Purchase Type	0.000000
Payment Method	0.000000
Railcard	0.000000
Ticket Class	0.000000
Ticket Type	0.000000
Price	0.000000
Departure Station	0.000000
Arrival Destination	0.000000
Date of Journey	0.000000
Departure Time	0.000000
Arrival Time	0.000000
Actual Arrival Time	5.939405
Journey Status	0.000000
Reason for Delay	0.000000
Refund Request	0.000000

dtype: float64

```
df['Actual Arrival Time'].fillna(df['Arrival Time'])
print(f"Missing values in 'Actual Arrival Time' after handling: {df['Actual Arrival Time'].isnull().sum()}")
```

```
➡ Missing values in 'Actual Arrival Time' after handling: 1880
```

```
print(df.isnull().mean() * 100)
```

```
➡ Transaction ID      0.0
   Date of Purchase   0.0
   Time of Purchase   0.0
   Purchase Type      0.0
   Payment Method     0.0
   Railcard           0.0
   Ticket Class       0.0
   Ticket Type        0.0
   Price             0.0
   Departure Station  0.0
   Arrival Destination 0.0
   Date of Journey    0.0
   Departure Time     0.0
   Arrival Time       0.0
   Actual Arrival Time 0.0
   Journey Status     0.0
   Reason for Delay   0.0
   Refund Request     0.0
dtype: float64
```

## ✓ Handle duplicates

```
df.duplicated().sum()
```

```
➡ np.int64(0)
```

## ✓ Data Transformation

```
df['Date of Purchase'] = pd.to_datetime(df['Date of Purchase'])
```

```
df['Date of Journey'] = pd.to_datetime(df['Date of Journey'])
```

```
df['Time of Purchase'] = pd.to_datetime(df['Time of Purchase'], format='%H:%M:%S')
```

```
df['Departure Time'] = pd.to_datetime(df['Departure Time'], format='%H:%M:%S')
```

```
df['Arrival Time'] = pd.to_datetime(df['Arrival Time'], format='%H:%M:%S')
```

```
df['Actual Arrival Time'] = pd.to_datetime(df['Actual Arrival Time'], format='%H:%M:%S', errors='coerce')
```

```
df.info()
```

```
➡ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 31653 entries, 0 to 31652
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Transaction ID         31653 non-null  object
1   Date of Purchase       31653 non-null  datetime64[ns]
2   Time of Purchase       31653 non-null  datetime64[ns]
3   Purchase Type          31653 non-null  object
4   Payment Method         31653 non-null  object
5   Railcard               31653 non-null  object
6   Ticket Class           31653 non-null  object
7   Ticket Type            31653 non-null  object
8   Price                  31653 non-null  int64
9   Departure Station      31653 non-null  object
10  Arrival Destination     31653 non-null  object
11  Date of Journey         31653 non-null  datetime64[ns]
12  Departure Time          31653 non-null  datetime64[ns]
13  Arrival Time            31653 non-null  datetime64[ns]
14  Actual Arrival Time     29773 non-null  datetime64[ns]
15  Journey Status         31653 non-null  object
16  Reason for Delay        31653 non-null  object
17  Refund Request          31653 non-null  object
dtypes: datetime64[ns](6), int64(1), object(11)
memory usage: 4.3+ MB
```

## ✓ Calculate journey duration and delay and convert categorical columns to category data type.

```
df['Departure_datetime'] = pd.to_datetime(df['Date of Journey'].astype(str) + ' ' + df['Departure Time'].astype(str))
```

```
df['Scheduled_Arrival_datetime'] = pd.to_datetime(df['Date of Journey'].astype(str) + ' ' + df['Arrival Time'].astype(str))
```

```
# Handle potential NaT values in 'Actual Arrival Time' before combining with 'Date of Journey'
df['Actual_Arrival_datetime'] = pd.to_datetime(df['Date of Journey'].astype(str) + ' ' + df['Actual Arrival Time'].astype(str), errors=

df['Journey Duration'] = (df['Actual_Arrival_datetime'] - df['Departure_datetime']).dt.total_seconds() / 60
df.loc[df['Journey Duration'] < 0, 'Journey Duration'] += 24 * 60 # Add 24 hours if negative

df['Delay'] = (df['Scheduled_Arrival_datetime'] - df['Actual_Arrival_datetime']).dt.total_seconds() / 60

categorical_cols = ['Purchase Type', 'Payment Method', 'Railcard', 'Ticket Class', 'Ticket Type', 'Departure Station', 'Arrival Destination']
for col in categorical_cols:
    df[col] = df[col].astype('category')

df = df.drop(columns=['Departure_datetime', 'Actual_Arrival_datetime', 'Scheduled_Arrival_datetime'])

⚡ /tmp/ipython-input-266334473.py:1: UserWarning: Could not infer format, so each element will be parsed individually, falling back to
df['Departure_datetime'] = pd.to_datetime(df['Date of Journey'].astype(str) + ' ' + df['Departure Time'].astype(str))
/tmp/ipython-input-266334473.py:2: UserWarning: Could not infer format, so each element will be parsed individually, falling back to
df['Scheduled_Arrival_datetime'] = pd.to_datetime(df['Date of Journey'].astype(str) + ' ' + df['Arrival Time'].astype(str))
/tmp/ipython-input-266334473.py:5: UserWarning: Could not infer format, so each element will be parsed individually, falling back to
df['Actual_Arrival_datetime'] = pd.to_datetime(df['Date of Journey'].astype(str) + ' ' + df['Actual Arrival Time'].astype(str), er
```

## ✓ Handle potential inconsistencies in categorical data

```
categorical_cols = ['Purchase Type', 'Payment Method', 'Railcard', 'Ticket Class', 'Ticket Type', 'Departure Station', 'Arrival Destination']
for col in categorical_cols:
    encoded_cols = [c for c in df.columns if c.startswith(col + '_')]
    print(f"Unique values for one-hot encoded '{col}' columns:")
    for encoded_col in encoded_cols:
        print(f"- {encoded_col}: {df[encoded_col].unique()}")
```

```
⚡ Unique values for one-hot encoded 'Purchase Type' columns:
Unique values for one-hot encoded 'Payment Method' columns:
Unique values for one-hot encoded 'Railcard' columns:
Unique values for one-hot encoded 'Ticket Class' columns:
Unique values for one-hot encoded 'Ticket Type' columns:
Unique values for one-hot encoded 'Departure Station' columns:
Unique values for one-hot encoded 'Arrival Destination' columns:
Unique values for one-hot encoded 'Journey Status' columns:
Unique values for one-hot encoded 'Reason for Delay' columns:
Unique values for one-hot encoded 'Refund Request' columns:
```

Based on the previous output, there are inconsistencies in the 'Reason for Delay' column with both 'Signal Failure' and 'Signal failure', and 'Staff Shortage' and 'Staffing', and 'Weather' and 'Weather Conditions'. These need to be unified by renaming the one-hot encoded columns and then combining them.

```
# Perform one-hot encoding on the 'Reason for Delay' column
df = pd.get_dummies(df, columns=['Reason for Delay'], prefix='Reason for Delay')

# Now combine the columns with similar meanings
df['Reason for Delay_Signal Failure'] = df['Reason for Delay_Signal Failure'] | df['Reason for Delay_Signal failure']
df.drop(columns=['Reason for Delay_Signal failure'], inplace=True)
df['Reason for Delay_Staff Shortage'] = df['Reason for Delay_Staff Shortage'] | df['Reason for Delay_Staffing']
df.drop(columns=['Reason for Delay_Staffing'], inplace=True)
df['Reason for Delay_Weather'] = df['Reason for Delay_Weather'] | df['Reason for Delay_Weather Conditions']
df.drop(columns=['Reason for Delay_Weather Conditions'], inplace=True)

# Display the updated columns
display(df[['Reason for Delay_Signal Failure', 'Reason for Delay_Staff Shortage', 'Reason for Delay_Weather']].head())
```

```
⚡
```


	Reason for Delay_Signal Failure	Reason for Delay_Staff Shortage	Reason for Delay_Weather
0	False	False	False
1	True	False	False
2	False	False	False
3	False	False	False
4	False	False	False

## Part 3

### Basic Analysis (Distribution , Outliers , Winsorizing)

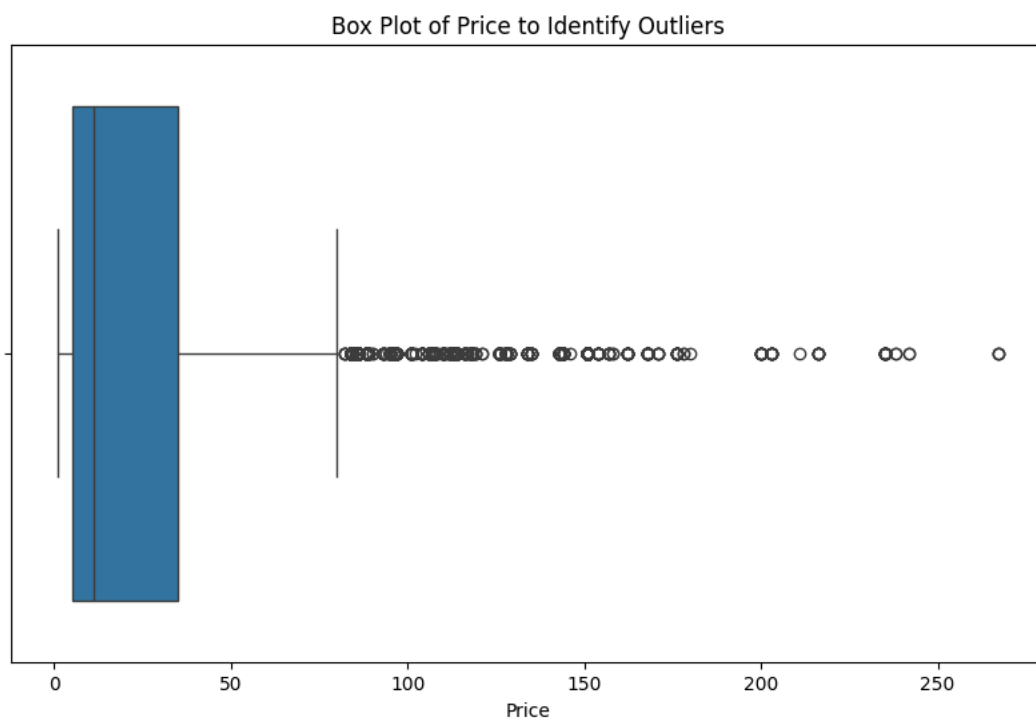
## ▼ Price Distribution

```
plt.figure(figsize=(10, 6))
sns.histplot(df['Price'], kde=True)
plt.title('Distribution of Price')
plt.xlabel('Price')
plt.ylabel('Frequency')
plt.show()
```



## ▼ Price Distribution (Outliers)

```
plt.figure(figsize=(10, 6))
sns.boxplot(x=df['Price'])
plt.title('Box Plot of Price to Identify Outliers')
plt.xlabel('Price')
plt.show()
```



Explore Outliers in Price

```
# Sort the DataFrame by 'Price' in descending order and display the top N rows
n_top_outliers = 10 # You can adjust this number
df_sorted_by_price = df.sort_values(by='Price', ascending=False)
display(df_sorted_by_price.head(n_top_outliers))
```

↔

	Transaction ID	Date of Purchase	Time of Purchase	Purchase Type	Payment Method	Railcard	Ticket Class	Ticket Type	Price	Departure Station	...	Actual Arrival Time	Journey Status	R Re
11849	de723682-3979-4d69-9664	2024-02-09	1900-01-01 06:30:59	Station	Credit Card	Unknown	First Class	Anytime	267	Manchester Piccadilly	...	1900-01-01 10:16:00	Delayed	
711	092e5598-08de-42f5-b6b2	2024-01-04	1900-01-01 06:35:01	Station	Credit Card	Unknown	First Class	Anytime	267	Manchester Piccadilly	...	1900-01-01 10:23:00	Delayed	
2042	05193f47-2107-4bb8-8adb	2024-01-09	1900-01-01 06:33:37	Station	Credit Card	Unknown	First Class	Anytime	267	Manchester Piccadilly	...	1900-01-01 10:29:00	Delayed	
31434	a4bbac34-6ed7-4d71-b738	2024-04-29	1900-01-01 16:49:20	Station	Contactless	Unknown	First Class	Anytime	242	Reading	...	1900-01-01 20:45:00	On Time	
13367	d327e0ec-e1ac-436a-aa5c	2024-02-13	1900-01-01 16:51:59	Station	Contactless	Unknown	First Class	Anytime	242	Reading	...	1900-01-01 20:45:00	On Time	
22488	419bcd59-fbc0-48f5-b48c	2024-03-26	1900-01-01 07:29:28	Station	Contactless	Unknown	First Class	Anytime	238	Liverpool Lime Street	...	1900-01-01 11:15:00	On Time	
20279	8e02cccf-5f58-48ca-b8b8	2024-03-18	1900-01-01 07:25:54	Station	Contactless	Unknown	First Class	Anytime	238	Liverpool Lime Street	...	1900-01-01 11:15:00	On Time	
24788	520bc09b-f83f-404a-9bf5	2024-04-05	1900-01-01 06:32:18	Station	Credit Card	Unknown	First Class	Anytime	235	Liverpool Lime Street	...	1900-01-01 11:11:00	Delayed	
21094	82b3b685-a3d7-49d1-8ffb	2024-03-21	1900-01-01 06:38:11	Station	Credit Card	Unknown	First Class	Anytime	235	Liverpool Lime Street	...	1900-01-01 10:48:00	Delayed	
29422	a94dd34b-3925-4eab-8b27	2024-04-22	1900-01-01 06:37:31	Station	Credit Card	Unknown	First Class	Anytime	235	Liverpool Lime Street	...	NaT	Cancelled	

10 rows × 24 columns

Define a threshold for high prices:

Calculate Q1, Q3, and IQR for the 'Price' column, define the upper bound for outliers, and use it as the threshold for high prices.

```
Q1 = df['Price'].quantile(0.25)
Q3 = df['Price'].quantile(0.75)
IQR = Q3 - Q1
upper_bound = Q3 + 1.5 * IQR
print(f"Q1: {Q1}")
print(f"Q3: {Q3}")
print(f"IQR: {IQR}")
print(f"Upper bound (High Price Threshold): {upper_bound}")
```

↔

```
Q1: 5.0
Q3: 35.0
IQR: 30.0
Upper bound (High Price Threshold): 80.0
```

Create a subset of high-priced tickets

Filter the DataFrame to create a new DataFrame containing only the rows where the 'Price' is above the defined threshold.

```
df_high_price = df[df['Price'] > upper_bound].copy()
display(df_high_price.head())
```



```
print(f"Number of high-priced tickets: {len(df_high_price)}")
```

↩

	Transaction ID	Date of Purchase	Time of Purchase	Purchase Type	Payment Method	Railcard	Ticket Class	Ticket Type	Price	Departure Station	...	Actual Arrival Time	Journey Status	Refi Requ
25	842da93c-b820-42dc-ad4f	2023-12-31	1900-01-01 15:19:53	Online	Contactless	Unknown	Standard	Advance	86	Manchester Piccadilly	...	1900-01-01 16:00:00	On Time	
45	767314a0-f839-4607-a3d3	2024-01-01	1900-01-01 05:09:30	Station	Credit Card	Unknown	First Class	Advance	134	Manchester Piccadilly	...	1900-01-01 05:31:00	Delayed	
51	382d60f9-9fe0-4920-97e4	2024-01-01	1900-01-01 06:34:08	Station	Credit Card	Unknown	Standard	Anytime	151	Liverpool Lime Street	...	1900-01-01 10:39:00	Delayed	
61	711c08ba-eb61-44ba-821a	2024-01-01	1900-01-01 09:30:09	Station	Credit Card	Unknown	First Class	Advance	134	Manchester Piccadilly	...	1900-01-01 10:08:00	Delayed	
68	9082a416-480e-4ca4-bf9d	2024-01-01	1900-01-01 15:39:11	Station	Credit Card	Unknown	Standard	Anytime	151	Liverpool Lime Street	...	1900-01-01 19:15:00	On Time	

5 rows × 24 columns

Number of high-priced tickets: 1555

▼ Compare characteristics

Create a subset for non-high-priced tickets, calculate and print descriptive statistics for numerical columns, and calculate and print value counts for categorical columns for both subsets.

```
df_not_high_price = df[df['Price'] <= upper_bound].copy()
numerical_cols_compare = ['Journey Duration', 'Delay']
print("Descriptive statistics for High-Priced Tickets:")
display(df_high_price[numerical_cols_compare].describe())
print("\nDescriptive statistics for Not High-Priced Tickets:")
display(df_not_high_price[numerical_cols_compare].describe())
categorical_cols_compare = ['Purchase Type', 'Payment Method', 'Railcard', 'Ticket Class', 'Ticket Type', 'Departure Station', 'Arrival Time']
print("\nValue counts for Categorical Columns (High-Priced Tickets):")
for col in categorical_cols_compare:
    print(f"\n{col}:")
    display(df_high_price[col].value_counts(normalize=True))
print("\nValue counts for Categorical Columns (Not High-Priced Tickets):")
for col in categorical_cols_compare:
    print(f"\n{col}:")
    display(df_not_high_price[col].value_counts(normalize=True))
```

↩ Descriptive statistics for High-Priced Tickets:

-----

KeyError

Traceback (most recent call last)

/tmp/ipython-input-332978330.py in <cell line: 0>()

2 numerical\_cols\_compare = ['Journey Duration', 'Delay']

3 print("Descriptive statistics for High-Priced Tickets:")

----> 4 display(df\_high\_price[numerical\_cols\_compare].describe())

5 print("\nDescriptive statistics for Not High-Priced Tickets:")

6 display(df\_not\_high\_price[numerical\_cols\_compare].describe())

-----

2 frames

/usr/local/lib/python3.11/dist-packages/pandas/core/indexes/base.py in \_raise\_if\_missing(self, key, indexer, axis\_name)

6247 if nmissing:

6248 if nmissing == len(indexer):

-> 6249 raise KeyError(f"None of [{key}] are in the [{axis\_name}]")

6250

6251 not\_found = list(ensure\_index(key)[missing\_mask.nonzero()[0]].unique())

KeyError: "None of [Index(['Journey Duration', 'Delay'], dtype='object')] are in the [columns]"

Visualize the distribution of key numerical and categorical columns for both high-priced and not high-priced tickets to compare their characteristics.

```
numerical_cols_to_plot = ['Journey Duration', 'Delay']
for col in numerical_cols_to_plot:
    plt.figure(figsize=(12, 6))
    sns.histplot(df_high_price[col], kde=True, color='skyblue', label='High Price')
    sns.histplot(df_not_high_price[col], kde=True, color='salmon', label='Not High Price')
    plt.title(f'Distribution of {col} by Price Category')
```


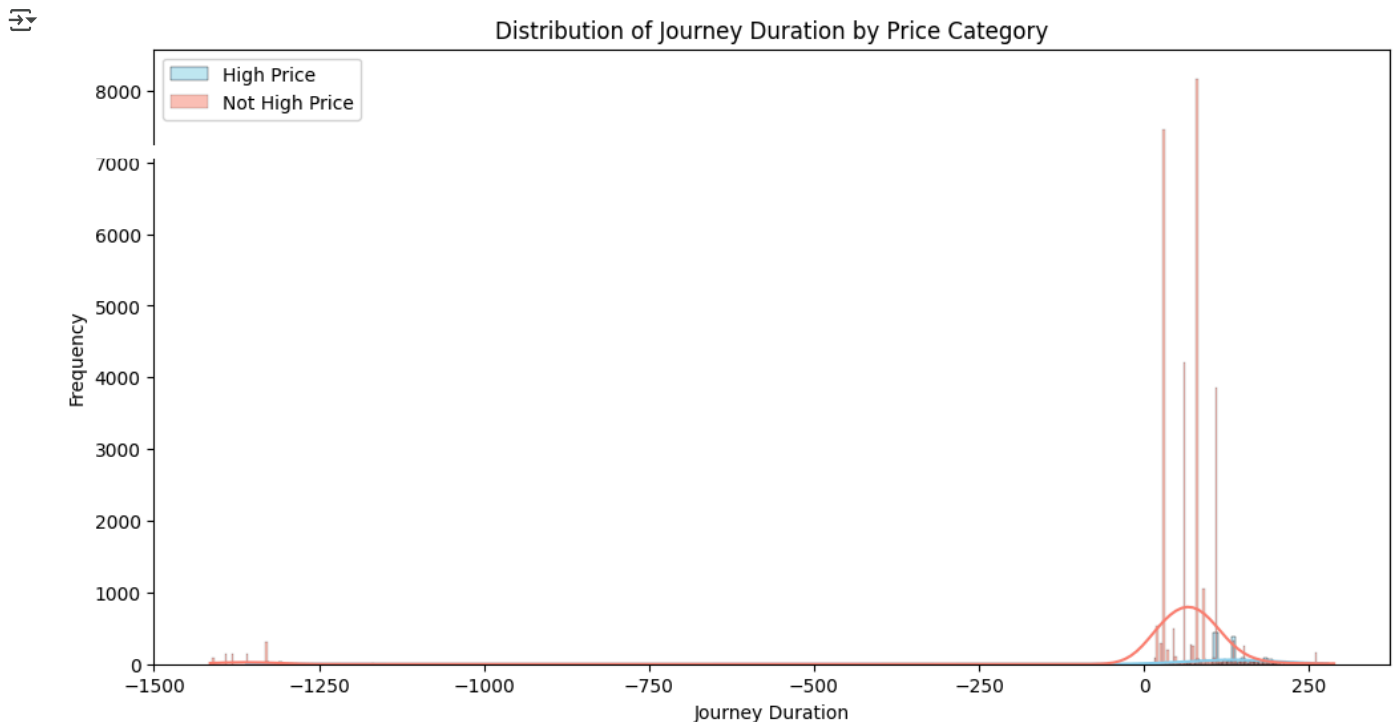
```

plt.xlabel(col)
plt.ylabel('Frequency')
plt.legend()
plt.show()

# Perform one-hot encoding on the categorical columns to be plotted
categorical_cols_to_encode = ['Purchase Type', 'Ticket Class', 'Ticket Type', 'Journey Status', 'Refund Request']
df_high_price_encoded = pd.get_dummies(df_high_price, columns=categorical_cols_to_encode, drop_first=False)
df_not_high_price_encoded = pd.get_dummies(df_not_high_price, columns=categorical_cols_to_encode, drop_first=False)

categorical_cols_to_plot = ['Purchase Type_Station', 'Ticket Class_Standard', 'Ticket Type_Anytime', 'Ticket Type_Off-Peak', 'Journey S
for col in categorical_cols_to_plot:
    plt.figure(figsize=(8, 5))
    high_price_counts = df_high_price_encoded[col].value_counts(normalize=True).reset_index()
    high_price_counts['Price Category'] = 'High Price'
    not_high_price_counts = df_not_high_price_encoded[col].value_counts(normalize=True).reset_index()
    not_high_price_counts['Price Category'] = 'Not High Price'
    combined_counts = pd.concat([high_price_counts, not_high_price_counts])
    sns.barplot(x=col, y='proportion', hue='Price Category', data=combined_counts)
    plt.title(f'Proportion of {col} by Price Category')
    plt.xlabel(col)
    plt.ylabel('Proportion')
    plt.show()

```

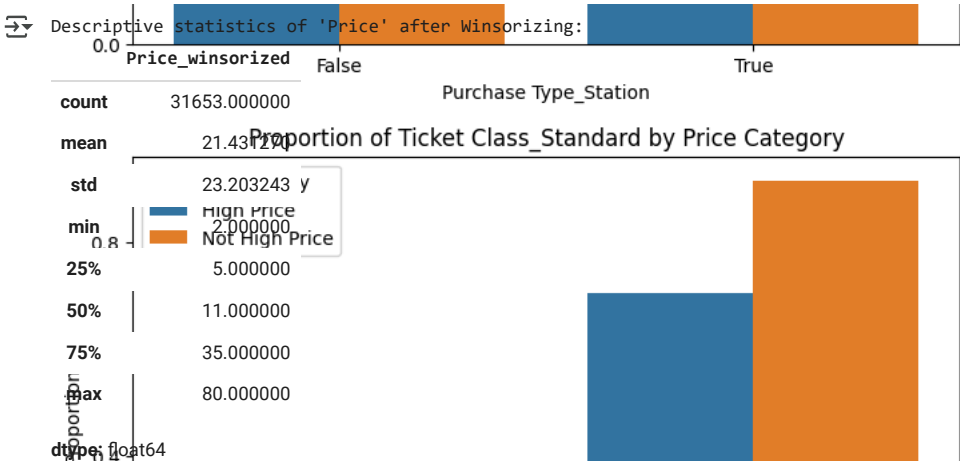


Apply Winsorizing to the 'Price' column to cap outlier values at specified percentiles.


Proportion of Purchase Type\_Station by Price Category

Calculate the lower and upper bounds for Winsorizing based on percentiles (e.g., 5th and 95th) and apply the capping to the 'Price' column. Display the descriptive statistics and a box plot after Winsorizing to verify the effect.

```
# Calculate lower and upper bounds based on percentiles
lower_bound = df['Price'].quantile(0.05)
upper_bound_winsorize = df['Price'].quantile(0.95)
# Apply Winsorizing
df['Price_winsorized'] = df['Price'].clip(lower=lower_bound, upper=upper_bound_winsorize)
# Display descriptive statistics of the new column
print("Descriptive statistics of 'Price' after Winsorizing:")
display(df['Price_winsorized'].describe())
# Display a box plot of the new column to visualize the effect
plt.figure(figsize=(10, 6))
sns.boxplot(x=df['Price_winsorized'])
plt.title('Box Plot of Price after Winsorizing')
plt.xlabel('Price (Winsorized)')
plt.show()
```



Box Plot of Price after Winsorizing




Part 4



KPIs

Proportion of Ticket Type\_Off-Peak by Price Category




```
# Display previously calculated KPIs for Revenue/Prices
print("Total Revenue:")
total_revenue = df['Price'].sum()
```

Key Performance Indicators: Revenue/Prices and Routes

```
print(f"{total_revenue}")


print("\nAverage Ticket Price:")
average_price = df['Price'].mean()
print(f"{average_price:.2f}")
```



## Revenue by Stations

```
# Calculate total revenue per departure station
revenue_by_departure_station = df.groupby('Departure Station')['Price'].sum().sort_values(ascending=False)
print("Total Revenue by Departure Station:")
display(revenue_by_departure_station)

# Calculate total revenue per arrival station
revenue_by_arrival_station = df.groupby('Arrival Destination')['Price'].sum().sort_values(ascending=False)
print("\nTotal Revenue by Arrival Station:")
display(revenue_by_arrival_station)
```



## Revenue by Ticket Type and Class

```
# Calculate total revenue by Ticket Type
revenue_by_ticket_type = df.groupby('Ticket Type')['Price'].sum().sort_values(ascending=False)
print("Total Revenue by Ticket Type:")
display(revenue_by_ticket_type)

# Calculate total revenue by Ticket Class
revenue_by_ticket_class = df.groupby('Ticket Class')['Price'].sum().sort_values(ascending=False)
print("\nTotal Revenue by Ticket Class:")
display(revenue_by_ticket_class)

# Calculate total revenue by Ticket Type and Ticket Class
revenue_by_type_and_class = df.groupby(['Ticket Type', 'Ticket Class'])['Price'].sum().unstack(fill_value=0)
print("\nTotal Revenue by Ticket Type and Ticket Class:")
display(revenue_by_type_and_class)
```



## Revenue per Journey

```
# Calculate average revenue per journey (which is the average ticket price)
average_revenue_per_journey = df['Price'].mean()
print(f"Average Revenue per Journey: {average_revenue_per_journey:.2f}")
```



## Key Performance Indicators: Time and Delay

```
# Calculate average delay
average_delay = df['Delay'].mean()
print(f"Average Delay (minutes): {average_delay:.2f}")

# Calculate proportion of delayed journeys
delayed_journeys_count = df[df['Journey Status'] == 'Delayed'].shape[0]
total_journeys_count = df.shape[0]
proportion_delayed = delayed_journeys_count / total_journeys_count
print(f"Proportion of Delayed Journeys: {proportion_delayed:.2%}")

# Calculate average delay for delayed journeys
average_delay_for_delayed = df[df['Journey Status'] == 'Delayed']['Delay'].mean()
print(f"Average Delay for Delayed Journeys Only (minutes): {average_delay_for_delayed:.2f}")
```

```
Average Delay (minutes): 3.06
Proportion of Delayed Journeys: 7.24%
Average Delay for Delayed Journeys Only (minutes): 42.21
```

## Key Performance Indicators: Ticket Types and Purchase Methods

```
# Calculate proportion of each Ticket Type
print("\nProportion of each Ticket Type:")
display(df['Ticket Type'].value_counts(normalize=True))

# Calculate proportion of each Purchase Type
print("\nProportion of each Purchase Type:")
display(df['Purchase Type'].value_counts(normalize=True))

# Calculate proportion of each Payment Method
print("\nProportion of each Payment Method:")
display(df['Payment Method'].value_counts(normalize=True))
```



Proportion of each Ticket Type:

proportion	
Ticket Type	
Advance	0.554797
Off-Peak	0.276498
Anytime	0.168704

**dtype:** float64

Proportion of each Purchase Type:

proportion	
Purchase Type	
Online	0.585126
Station	0.414874

**dtype:** float64

Proportion of each Payment Method:

proportion	
Payment Method	
Credit Card	0.604556
Contactless	0.342274
Debit Card	0.053170

**dtype:** float64

## ✓ Part 5

## Deep analysis

### Analyze the distribution of the price column

```
df_not_high_price = df[df['Price'] <= upper_bound].copy()

numerical_cols_compare = ['Journey Duration', 'Delay']
print("Descriptive statistics for High-Priced Tickets:")
display(df_high_price[numerical_cols_compare].describe())
print("\nDescriptive statistics for Not High-Priced Tickets:")
display(df_not_high_price[numerical_cols_compare].describe())

categorical_cols_compare = ['Purchase Type', 'Payment Method', 'Railcard', 'Ticket Class', 'Ticket Type', 'Departure Station', 'Arrival Station']
print("\nValue counts for Categorical Columns (High-Priced Tickets):")
for col in categorical_cols_compare:
    print(f"\n{col}:")
    display(df_high_price[col].value_counts(normalize=True))

print("\nValue counts for Categorical Columns (Not High-Priced Tickets):")
for col in categorical_cols_compare:
    print(f"\n{col}:")
    display(df_not_high_price[col].value_counts(normalize=True))

numerical_cols_to_plot = ['Journey Duration', 'Delay']
for col in numerical_cols_to_plot:
    plt.figure(figsize=(12, 6))
    sns.histplot(df_high_price[col], kde=True, color='skyblue', label='High Price')
    sns.histplot(df_not_high_price[col], kde=True, color='salmon', label='Not High Price')
```

```
plt.title(f'Distribution of {col} by Price Category')
plt.xlabel(col)
plt.ylabel('Frequency')
plt.legend()
plt.show()

categorical_cols_to_encode = ['Purchase Type', 'Ticket Class', 'Ticket Type', 'Journey Status', 'Refund Request']
df_high_price_encoded = pd.get_dummies(df_high_price, columns=categorical_cols_to_encode, drop_first=False)
df_not_high_price_encoded = pd.get_dummies(df_not_high_price, columns=categorical_cols_to_encode, drop_first=False)

categorical_cols_to_plot = [col for col in df_high_price_encoded.columns if any(cat in col for cat in categorical_cols_to_encode)]
for col in categorical_cols_to_plot:
    if df_high_price_encoded[col].nunique() > 1: # Only plot if there are at least two categories
        plt.figure(figsize=(8, 5))
        high_price_counts = df_high_price_encoded[col].value_counts(normalize=True).reset_index()
        high_price_counts['Price Category'] = 'High Price'
        not_high_price_counts = df_not_high_price_encoded[col].value_counts(normalize=True).reset_index()
        not_high_price_counts['Price Category'] = 'Not High Price'
        combined_counts = pd.concat([high_price_counts, not_high_price_counts])
        sns.barplot(x=col, y='proportion', hue='Price Category', data=combined_counts)
        plt.title(f'Proportion of {col} by Price Category')
        plt.xlabel(col)
        plt.ylabel('Proportion')
        plt.show()
```

🔗 Descriptive statistics for High-Priced Tickets:

	Journey Duration	Delay
count	1555.000000	1555.000000
mean	83.315113	12.932476
std	271.703414	19.448773
min	-1360.000000	0.000000
25%	110.000000	0.000000
50%	135.000000	0.000000
75%	150.000000	24.000000
max	277.000000	60.000000

Descriptive statistics for Not High-Priced Tickets:

	Journey Duration	Delay
count	30098.000000	30098.000000
mean	27.762742	2.546448
std	81.666667	10.000000
min	-1415.000000	0.000000
25%	30.000000	0.000000
50%	80.000000	0.000000
75%	80.000000	0.000000
max	288.000000	180.000000

Value counts for Categorical Columns (High-Priced Tickets):

Purchase Type:

	proportion
Purchase Type	
Station	0.62701
Online	0.37299

dtype: float64

Payment Method:

	proportion
Payment Method	
Credit Card	0.701608
Contactless	0.176849

dtype: float64

✓ Analyze relationship between price and other variables

dtype: float64

Subtask: Railcard:

proportion

Explore the correlation or relationship between 'Price' and other numerical and categorical variables to understand which factors might influence the ticket price.

Unknown 0.865595

Reasoning: Calculate and display the correlation matrix for the numerical columns and create a heatmap to visualize the relationships.

Disabled 0.033441

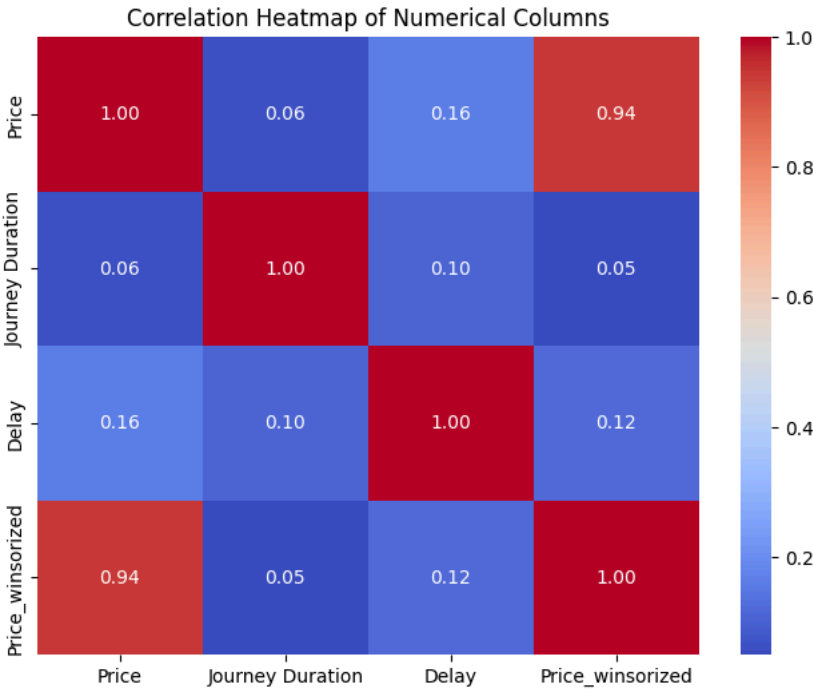
```
numerical_cols = ['Price', 'Journey Duration', 'Delay', 'Price_winsorized']
correlation_matrix = df[numerical_cols].corr()
print("Correlation Matrix for Numerical Columns:")
display(correlation_matrix)
```

```
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Heatmap of Numerical Columns')
plt.show()
```

Standard 0.104823

Correlation Matrix for Numerical Columns:

	First Class	Price	Journey Duration	Delay	Price_winsorized
Price	1.000000	0.295177	0.059552	0.158477	0.942523
Journey Duration	0.059552	1.000000	0.104464	0.104464	0.050069
Delay	0.158477	0.104464	1.000000	0.117862	0.117862
Price_winsorized	0.942523	0.050069	0.117862	1.000000	1.000000



Reasoning: Create box plots to visualize the distribution of 'Price' for key categorical columns and calculate the average 'Price' for each category.

dtune float64

```
categorical_cols_for_boxplot = ['Ticket Class', 'Ticket Type', 'Purchase Type', 'Journey Status', 'Payment Method']
for col in categorical_cols_for_boxplot:
    plt.figure(figsize=(10, 6))
    sns.boxplot(x=col, y='Price', data=df)
    plt.title(f'Price Distribution by {col}')
    plt.xlabel(col)
    plt.ylabel('Price')
    plt.show()
```

```
print(f"\nAverage Price by {col}:")
average_price_by_category = df.groupby(col, observed=True)['Price'].mean()
display(average_price_by_category)
```

Birmingham New Street 0.000000

Liverpool Lime Street 0.030868

London St Pancras 0.028939

London Kings Cross 0.008360

Peterborough 0.005145

Edinburgh Waverlev 0.002572




```

Stafford      0.000000
Average Price by Ticket Class:
Tamworth      0.000000
Price
Wakefield     0.000000
Ticket Class  0.000000
...
First Class   48.855134
Wolverhampton 0.000000
Standard     20.721175

dtype: float64
dtype: float64

```



## Analyze purchase trends over time

Resample the DataFrame by day, week, and month to analyze purchase trends over different time periods and then plot the results.

```

daily_purchases = df.set_index('Date of Purchase').resample('D').size().reset_index(name='purchase_count')
weekly_purchases = df.set_index('Date of Purchase').resample('W').size().reset_index(name='purchase_count')
monthly_purchases = df.set_index('Date of Purchase').resample('M').size().reset_index(name='purchase_count')

```

```

plt.figure(figsize=(12, 6))
sns.lineplot(x='Date of Purchase', y='purchase_count', data=daily_purchases)
plt.title('Daily Purchase Trends')
plt.xlabel('Date')
plt.ylabel('Number of Purchases')
plt.show()

```

```

plt.figure(figsize=(12, 6))
sns.lineplot(x='Date of Purchase', y='purchase_count', data=weekly_purchases)
plt.title('Weekly Purchase Trends')
plt.xlabel('Date')
plt.ylabel('Number of Purchases')
plt.show()

```

```



plt.figure(figsize=(12, 6))
sns.lineplot(x='Date of Purchase', y='purchase_count', data=monthly_purchases)
plt.title('Monthly Purchase Trends')
plt.xlabel('Date')

```



```
plt.ylabel('Number of Purchases')
plt.show()


FutureWarning: 'M' is deprecated and will be removed in a future version, please use 'ME' instead
monthly_purchases = df.set_index('Date of Purchase').resample('M').size().reset_index(name='purchase_count')
```

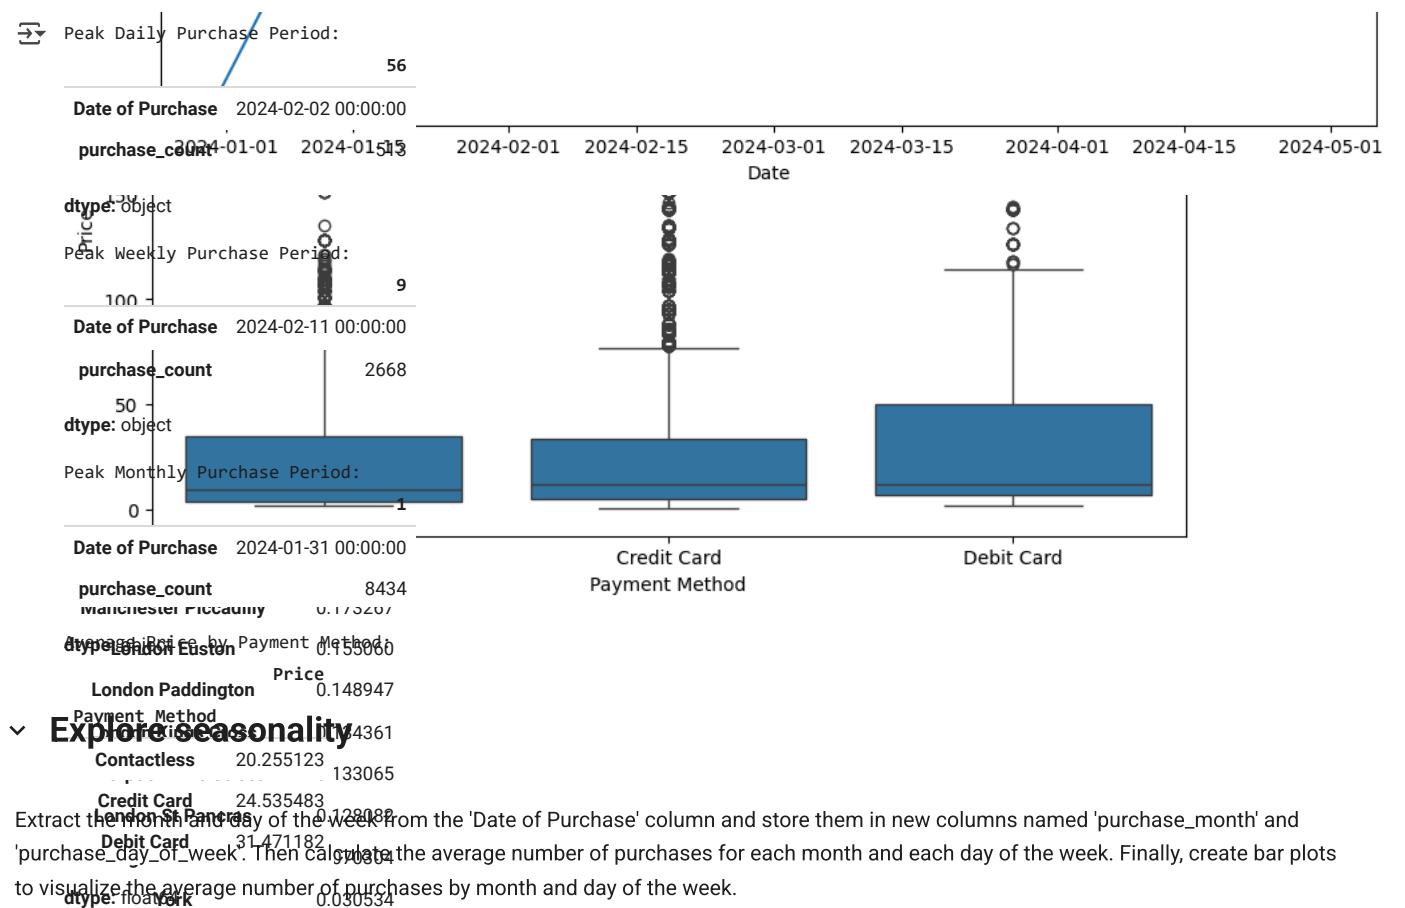


Identify peak purchase periods

Reasoning: Find and print the peak daily, weekly, and monthly purchase periods and their corresponding purchase counts.

```
peak_day = daily_purchases.loc[daily_purchases['purchase_count'].idxmax()]
peak_week = weekly_purchases.loc[weekly_purchases['purchase_count'].idxmax()]
peak_month = monthly_purchases.loc[monthly_purchases['purchase_count'].idxmax()]
print("Peak Daily Purchase Period:")
display(peak_day)
print("\nPeak Weekly Purchase Period:")
display(peak_week)
print("\nPeak Monthly Purchase Period:")
display(peak_month)
```





```
df['purchase_month'] = df['Date of Purchase'].dt.month
df['purchase_day_of_week'] = df['Date of Purchase'].dt.day_name()



average_purchases_by_month = df.groupby('purchase_month').size().reset_index(name='average_purchases')
average_purchases_by_month['average_purchases'] = average_purchases_by_month['average_purchases'] / len(df['Date of Purchase'].dt.to_period('M'))

average_purchases_by_day_of_week = df.groupby('purchase_day_of_week').size().reset_index(name='average_purchases')
days_order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
average_purchases_by_day_of_week['purchase_day_of_week'] = pd.Categorical(average_purchases_by_day_of_week['purchase_day_of_week'], categories=days_order)
average_purchases_by_day_of_week = average_purchases_by_day_of_week.sort_values('purchase_day_of_week')
average_purchases_by_day_of_week['average_purchases'] = average_purchases_by_day_of_week['average_purchases'] / len(df['Date of Purchase'].dt.to_period('D'))

plt.figure(figsize=(10, 6))
sns.barplot(x='purchase_month', y='average_purchases', data=average_purchases_by_month)
plt.title('Average Number of Purchases by Month')
plt.xlabel('Month')
plt.ylabel('Average Number of Purchases')
plt.xticks(ticks=range(12), labels=['Jan', 'Feb', 'Mar', 'Apr', 'May', 'Jun', 'Jul', 'Aug', 'Sep', 'Oct', 'Nov', 'Dec'])
plt.show()

plt.figure(figsize=(10, 6))
sns.barplot(x='purchase_day_of_week', y='average_purchases', data=average_purchases_by_day_of_week)
plt.title('Average Number of Purchases by Day of the Week')
plt.xlabel('Day of the Week')
plt.ylabel('Average Number of Purchases')
plt.show()
```

Leicester	0.011197
Sheffield	0.009037
Durham	0.008572
Leeds	0.008472
Peterborough	0.007775
Swindon	0.007575
Tamworth	0.007542
Nuneaton	0.007276
Doncaster	0.007010
London Paddington	0.006579
Crewe	0.006412
Stafford	0.006313



## ✓ Analyze purchase day of the week

Calculate the number of purchases for each day of the week, sort the results, print the counts, and identify the day with the highest number of purchases.

```
Reason for Delay_Technical_Issue
False    0.979268
True     0.020732

purchases_by_day_of_week = df['purchase_day_of_week'].value_counts()
print("Number of purchases for each day of the week:")
display(purchases_by_day_of_week)
most_popular_day = purchases_by_day_of_week.idxmax()
print(f"\nThe day of the week with the most purchases is: {most_popular_day}")
```

Reason for Delay_Traffic	
False	0.990099
True	0.009901

dtype: float64

Reason for Delay\_Unknown Reason:

Reason for Delay_Unknown Reason		proportion
True		0.883979
False		0.116021

➡ Number of purchases for each day of the week:  
**dtype: float64**                      **count**

Reason for Delay\_Weather:

Reason for Delay_Weather	count	proportion
Sunday	4676	
Friday	4627	
Wednesday	4602	0.969267
True Saturday	4477	0.030733
Monday	4455	
Tuesday	4454	
Thursday	4362	

Refund Request

Refund Request	count	proportion
No	46734	0.96734
Yes	1366	0.03266


The day of the week with the most purchases is: Sunday

## ✓ Analyze purchase time of day


Distribution of Journey Duration by Price Category

Extract the hour from 'Time of Purchase', count purchases per hour, print counts, and find the hour with the most purchases.


```
df['purchase_hour'] = df['Time of Purchase'].dt.hour
purchases_by_hour = df['purchase_hour'].value_counts().sort_index()
print("Number of purchases for each hour of the day:")
display(purchases_by_hour)
most_popular_hour = purchases_by_hour.idxmax()
print(f"\nThe hour of the day with the most purchases is: {most_popular_hour}")
```





Distribution of Delay by Price Category





Proportion of Purchase Type\_Online by Price Category






**Reasoning:** Create a bar plot to visualize the number of purchases by hour, add title and labels, and display the plot.

```
plt.figure(figsize=(12, 6))
sns.barplot(x=purchases_by_hour.index, y=purchases_by_hour.values)
plt.title('Number of Purchases by Hour of the Day')
plt.xlabel('Hour of the Day')
plt.ylabel('Number of Purchases')
plt.xticks(rotation=0)
plt.show()
```

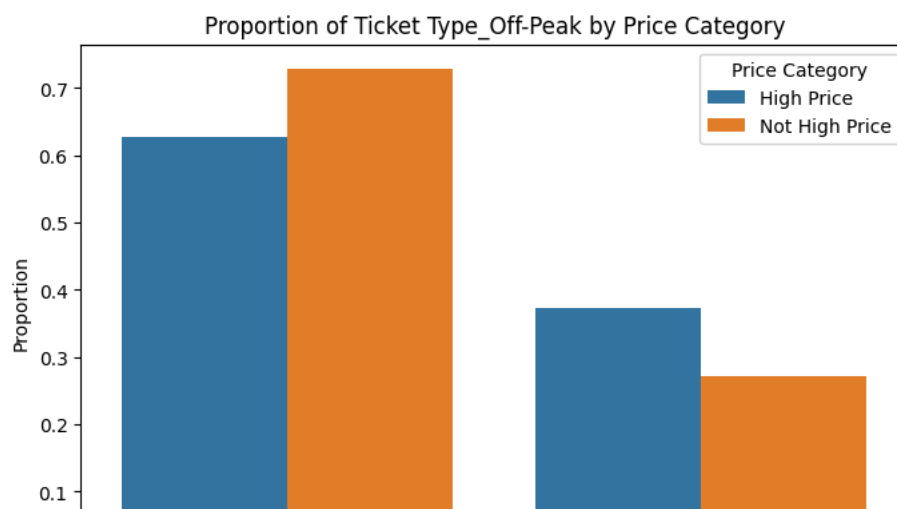
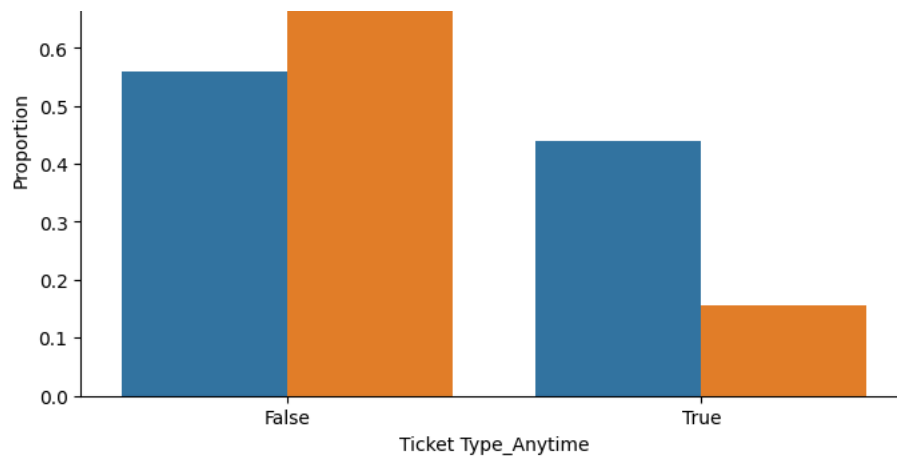


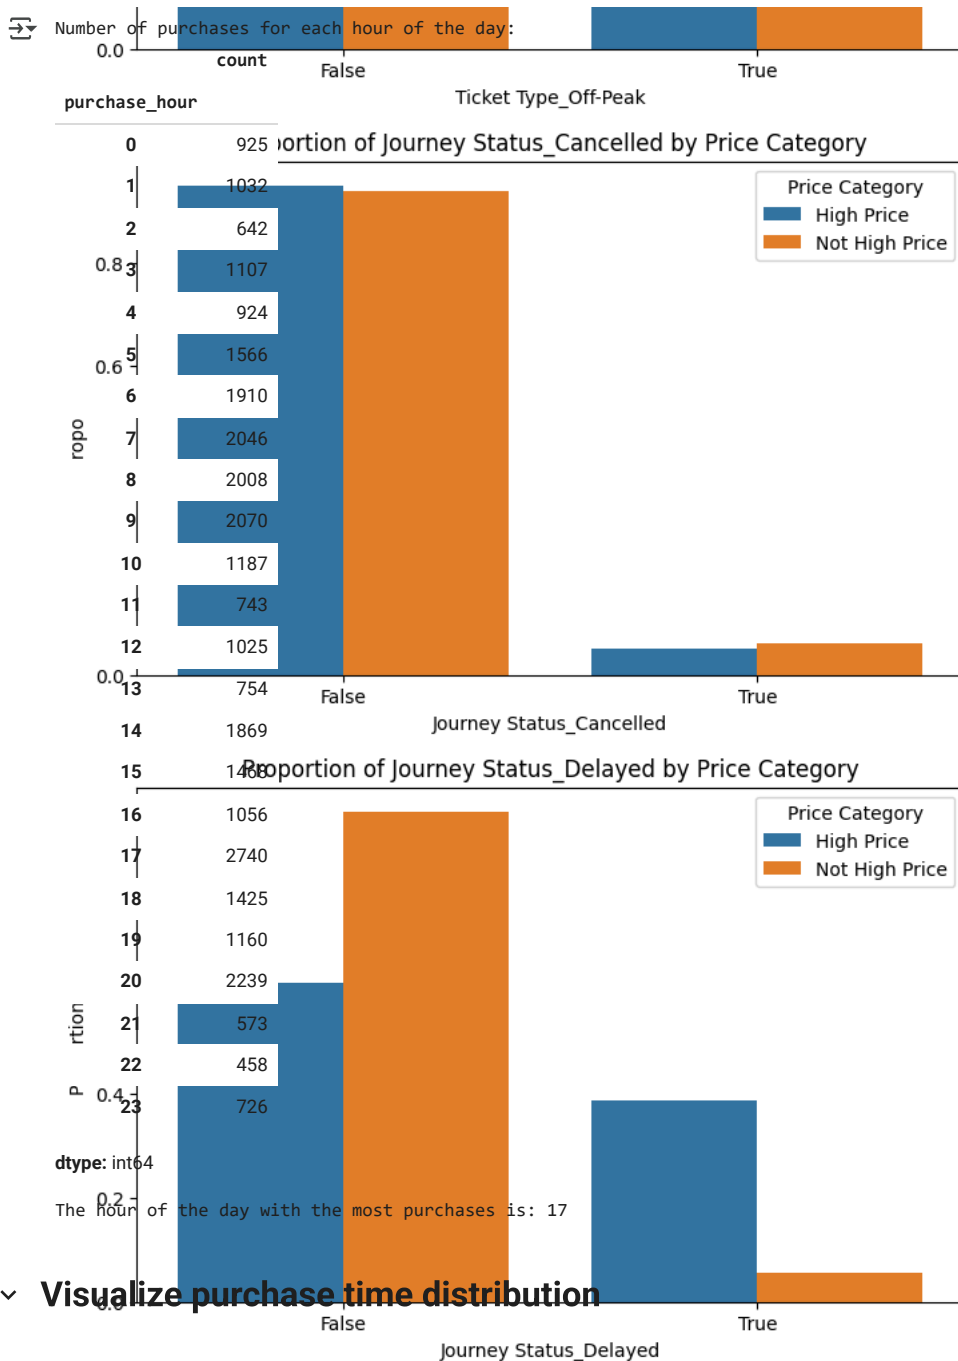


## ✓ Analyze purchase time of day

Extract the hour from 'Time of Purchase', count purchases per hour, print counts, and find the hour with the most purchases.


```
df['purchase_hour'] = df['Time of Purchase'].dt.hour
purchases_by_hour = df['purchase_hour'].value_counts().sort_index()
print("Number of purchases for each hour of the day:")
display(purchases_by_hour)
most_popular_hour = purchases_by_hour.idxmax()
print(f"\nThe hour of the day with the most purchases is: {most_popular_hour}")
```

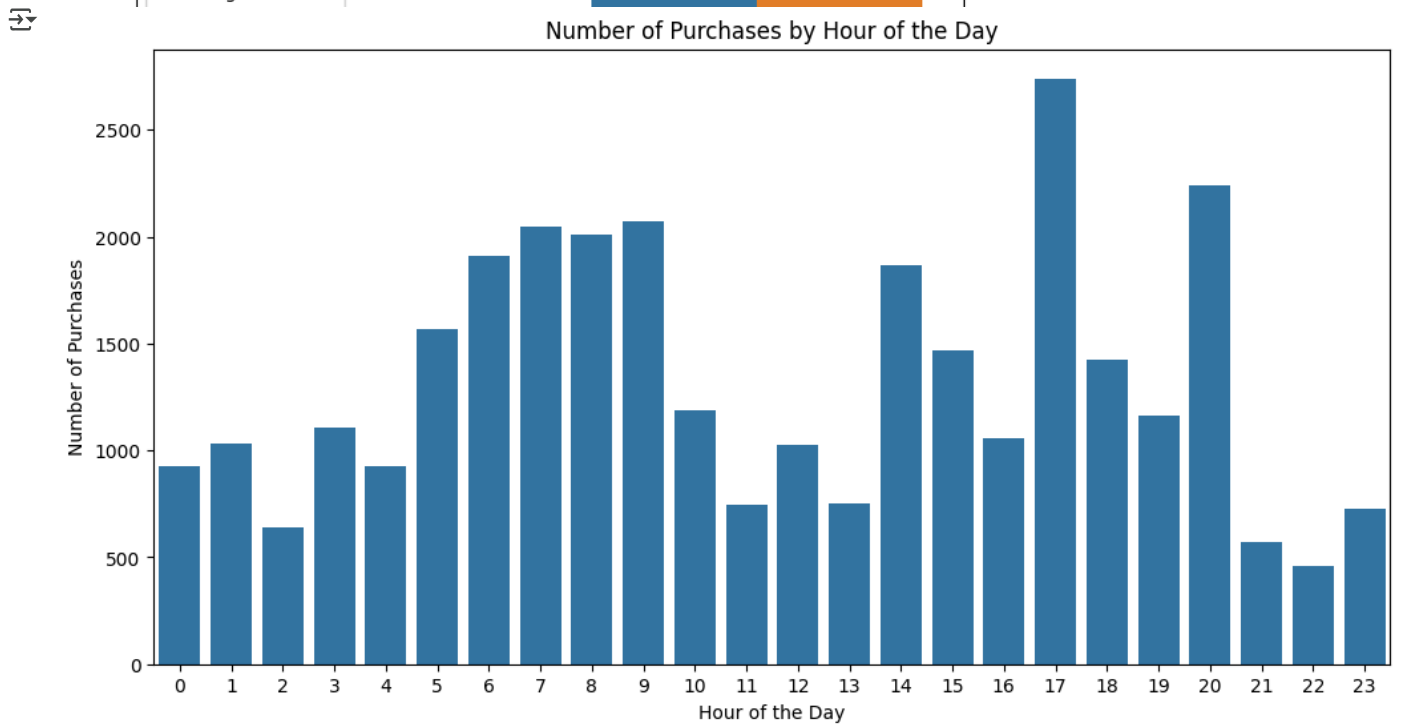




Create a bar plot to visualize the number of purchases by hour, add title and labels, and display the plot.

```
plt.figure(figsize=(12, 6))
sns.barplot(x=purchases_by_hour.index, y=purchases_by_hour.values)
plt.title('Number of Purchases by Hour of the Day')
plt.xlabel('Hour of the Day')
plt.ylabel('Number of Purchases')
plt.xticks(rotation=0)
plt.show()
```





## ✓ Analyze departure time distribution

Extract the hour from 'Departure Time', count occurrences per hour, sort by hour, print counts, find the hour with the most occurrences, and print the peak hour.

```
df['departure_hour'] = df['Departure Time'].dt.hour
departures_by_hour = df['departure_hour'].value_counts().sort_index()
print("Number of departures for each hour of the day:")
display(departures_by_hour)
most_popular_departure_hour = departures_by_hour.idxmax()
print(f"\nThe hour of the day with the most departures is: {most_popular_departure_hour}")
```



➡ Number of departures for each hour of the day:


count	
departure_hour	
0	853
1	644
2	942
3	543
4	1041
5	725
6	3112
7	2795
8	2179
9	1230
10	525
11	1143
12	773
13	1276
14	855
15	1220
16	2301
17	2888
18	3113
19	438
20	1058
21	570
22	788
23	641

**dtype:** int64

The hour of the day with the most departures is: 18

Visualize the number of departures by hour using a bar plot to better understand the distribution and visually confirm the peak hours.

```
plt.figure(figsize=(12, 6))
sns.barplot(x=departures_by_hour.index, y=departures_by_hour.values)
plt.title('Number of Departures by Hour of the Day')
plt.xlabel('Hour of the Day')
plt.ylabel('Number of Departures')
plt.xticks(rotation=0)
plt.show()
```



## ✓ Analyze arrival time distribution

Extract the hour from the 'Arrival Time' column, count the occurrences of each hour, sort the results, print the counts, and find the hour with the most arrivals.

```
df['arrival_hour'] = df['Arrival Time'].dt.hour
arrivals_by_hour = df['arrival_hour'].value_counts().sort_index()
print("Number of arrivals for each hour of the day:")
display(arrivals_by_hour)
most_popular_arrival_hour = arrivals_by_hour.idxmax()
print(f"\nThe hour of the day with the most arrivals is: {most_popular_arrival_hour}")
```

➡ Number of arrivals for each hour of the day:

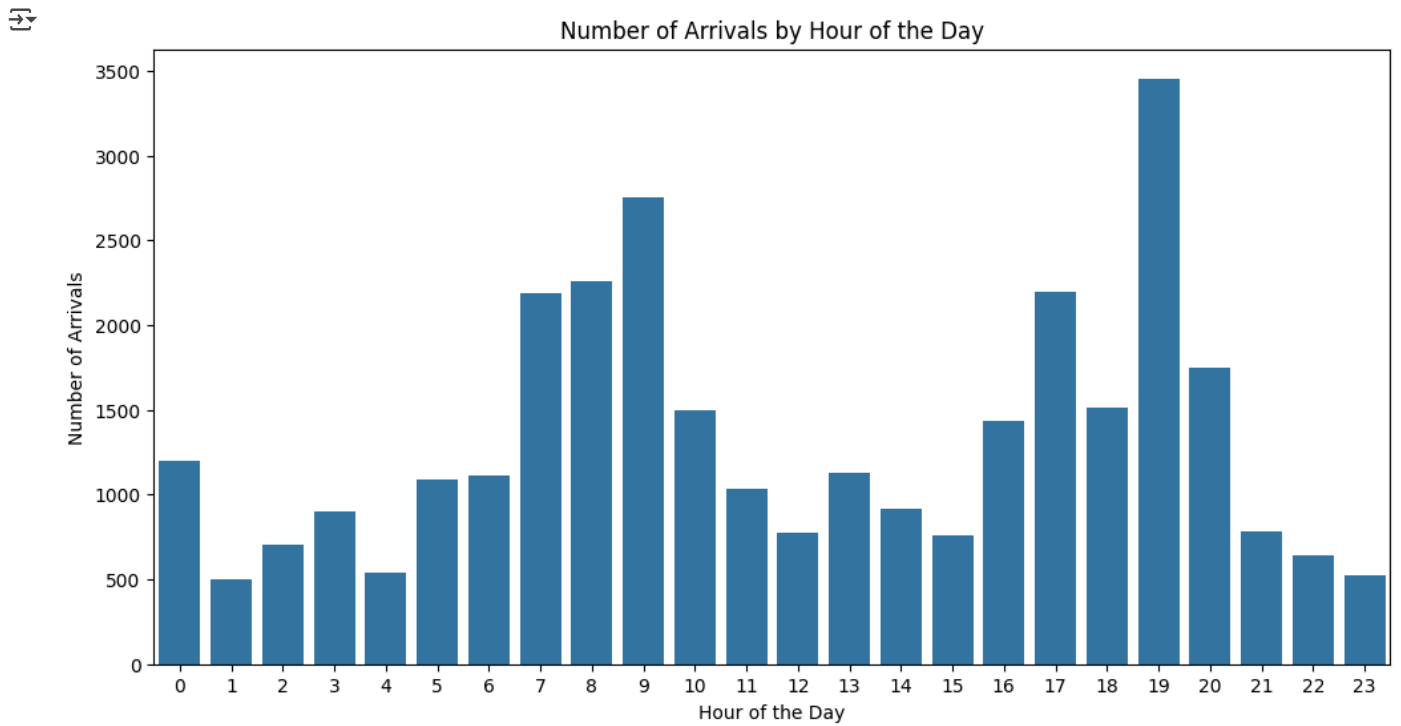
	count
arrival_hour	
0	1200
1	503
2	704
3	898
4	543
5	1090
6	1111
7	2185
8	2258
9	2752
10	1495
11	1037
12	775
13	1132
14	916
15	757
16	1436
17	2200
18	1510
19	3455
20	1751
21	780
22	645
23	520

**dtype:** int64

The hour of the day with the most arrivals is: 19

Create a bar plot to visualize the number of arrivals by hour, add title and labels, and display the plot.

```
plt.figure(figsize=(12, 6))
sns.barplot(x=arrivals_by_hour.index, y=arrivals_by_hour.values)
plt.title('Number of Arrivals by Hour of the Day')
plt.xlabel('Hour of the Day')
plt.ylabel('Number of Arrivals')
plt.xticks(rotation=0)
plt.show()
```



## ✓ Identify peak travel times

Based on the outputs of the previous subtasks, identify and display the peak travel periods.

```
print(f"Peak Departure Hour: {most_popular_departure_hour}")
print(f"Peak Arrival Hour: {most_popular_arrival_hour}")
print("\nBased on the analysis of departure and arrival times, the overall peak travel periods are in the early morning (around 7-9 AM), late
```

```
↔ Peak Departure Hour: 18
Peak Arrival Hour: 19
```

Based on the analysis of departure and arrival times, the overall peak travel periods are in the early morning (around 7-9 AM), late

## ✓ Geographical Analysis: Identifying Busiest Stations

```
# Calculate the number of departures from each station
departure_counts = df['Departure Station'].value_counts()
```

```
print("Number of departures from each station:")
display(departure_counts)
```

➡ Number of departures from each station:

count	
Departure Station	
Manchester Piccadilly	5650
London Euston	4954
Liverpool Lime Street	4561
London Paddington	4500
London Kings Cross	4229
London St Pancras	3891
Birmingham New Street	2136
York	927
Reading	594
Oxford	144
Edinburgh Waverley	51
Bristol Temple Meads	16

dtype: int64

```
# Calculate the number of arrivals at each station
arrival_counts = df['Arrival Destination'].value_counts()

print("\nNumber of arrivals at each station:")
display(arrival_counts)
```



Number of arrivals at each station:

	count
Arrival Destination	
Birmingham New Street	7742
Liverpool Lime Street	5022
York	4019
Manchester Piccadilly	3968
Reading	3920
London Euston	1567
London St Pancras	749
Oxford	623
London Paddington	351
Leicester	337
Sheffield	272
Durham	258
Leeds	255
Peterborough	242
Swindon	228
Tamworth	227
Nuneaton	219
Doncaster	211
Crewe	193
Stafford	190
Edinburgh Waverley	178
Nottingham	158
Edinburgh	154
Bristol Temple Meads	144
Wolverhampton	115
London Kings Cross	84
London Waterloo	68
Coventry	65
Didcot	48

Summary of Busiest Stations: 16

Based on the departure and arrival counts, we can identify the stations with the highest traffic.

Warrington 15

Wakefield 15

```
print("\nTop 5 Busiest Departure Stations:")
display(departure_counts.head())
```

```
print("\nTop 5 Busiest Arrival Stations:")
display(arrival_counts.head())
```



Top 5 Busiest Departure Stations:

	count
Departure Station	
Manchester Piccadilly	5650
London Euston	4954
Liverpool Lime Street	4561
London Paddington	4500
London Kings Cross	4229

dtype: int64

Top 5 Busiest Arrival Stations:

	count
Arrival Destination	
Birmingham New Street	7742
Liverpool Lime Street	5022
York	4019
Manchester Piccadilly	3968
Reading	3920


dtype: int64

## Visualize Important Routes

```
# Create a pivot table to show the number of journeys between departure and arrival stations
route_matrix = df.pivot_table(index='Departure Station', columns='Arrival Destination', values='Transaction ID', aggfunc='count', fill_value=0)

plt.figure(figsize=(14, 10))
sns.heatmap(route_matrix, annot=True, fmt='d', cmap='viridis')
plt.title('Network Heatmap of Journeys Between Stations')
plt.xlabel('Arrival Destination')
plt.ylabel('Departure Station')
plt.xticks(rotation=90)
plt.yticks(rotation=0)
plt.tight_layout()
plt.show()
```

```
/tmp/ipython-input-377728584.py:2: FutureWarning: The default value of observed=False is deprecated and will change to observed=True
route_matrix = df.pivot_table(index='Departure Station', columns='Arrival Destination', values='Transaction ID', aggfunc='count',
```



## Define routes

Concatenate the 'Departure Station' and 'Arrival Destination' columns to create the 'Route' column and display the head of the dataframe to confirm the creation of the new column.

```
df['Route'] = df['Departure Station'].astype(str) + ' to ' + df['Arrival Destination'].astype(str)
display(df.head())
```



↕

	Transaction ID	Date of Purchase	Time of Purchase	Purchase Type	Payment Method	Railcard	Ticket Class	Ticket Type	Price	Departure Station	...	Reason for Delay_Traffic	Reason for Delay_Unknown
0	da8a6ba8-b3dc-4677-b176	2023-12-08	1900-01-01 12:41:11	Online	Contactless	Adult	Standard	Advance	43	London Paddington	...	False	
1	b0cdd1b0-f214-4197-be53	2023-12-16	1900-01-01 11:23:01	Station	Credit Card	Adult	Standard	Advance	23	London Kings Cross	...	False	
2	f3ba7a96-f713-40d9-9629	2023-12-19	1900-01-01 19:51:27	Online	Credit Card	Unknown	Standard	Advance	3	Liverpool Lime Street	...	False	
3	b2471f11-4fe7-4c87-8ab4	2023-12-20	1900-01-01 23:00:36	Station	Credit Card	Unknown	Standard	Advance	13	London Paddington	...	False	
4	2be00b45-0762-485e-a7a3	2023-12-27	1900-01-01 18:22:56	Online	Contactless	Unknown	Standard	Advance	76	Liverpool Lime Street	...	False	

5 rows × 32 columns

The error indicates that the 'Departure Station' and 'Arrival Destination' columns are of categorical type and cannot be directly concatenated with a string. Convert these columns to string type before concatenation.

```
df['Route'] = df['Departure Station'].astype(str) + ' to ' + df['Arrival Destination'].astype(str)
display(df.head())
```

↕

	Transaction ID	Date of Purchase	Time of Purchase	Purchase Type	Payment Method	Railcard	Ticket Class	Ticket Type	Price	Departure Station	...	Reason for Delay_Traffic	Reason for Delay_Unknown
0	da8a6ba8-b3dc-4677-b176	2023-12-08	1900-01-01 12:41:11	Online	Contactless	Adult	Standard	Advance	43	London Paddington	...	False	
1	b0cdd1b0-f214-4197-be53	2023-12-16	1900-01-01 11:23:01	Station	Credit Card	Adult	Standard	Advance	23	London Kings Cross	...	False	
2	f3ba7a96-f713-40d9-9629	2023-12-19	1900-01-01 19:51:27	Online	Credit Card	Unknown	Standard	Advance	3	Liverpool Lime Street	...	False	
3	b2471f11-4fe7-4c87-8ab4	2023-12-20	1900-01-01 23:00:36	Station	Credit Card	Unknown	Standard	Advance	13	London Paddington	...	False	
4	2be00b45-0762-485e-a7a3	2023-12-27	1900-01-01 18:22:56	Online	Contactless	Unknown	Standard	Advance	76	Liverpool Lime Street	...	False	

5 rows × 32 columns

▼ Identify popular routes

Calculate and print the value counts for the 'Route' column to identify the most popular routes and find and print the route with the highest number of purchases.

```
route_counts = df['Route'].value_counts()
print("Number of purchases for each route:")
display(route_counts)
most_popular_route = route_counts.idxmax()
print(f"\nThe most popular route is: {most_popular_route}")
```

➡ Number of purchases for each route:

	count
Route	
Manchester Piccadilly to Liverpool Lime Street	4628
London Euston to Birmingham New Street	4209
London Kings Cross to York	3922
London Paddington to Reading	3873
London St Pancras to Birmingham New Street	3471
...	...
York to Edinburgh Waverley	15
York to Wakefield	15
York to Liverpool Lime Street	15
Manchester Piccadilly to Warrington	15
Liverpool Lime Street to Birmingham New Street	14

65 rows × 1 columns

dtype: int64

The most popular route is: Manchester Piccadilly to Liverpool Lime Street

## ✓ Analyze popular routes

Calculate the average price, journey duration, and delay for each route and print them sorted.

```
average_price_by_route = df.groupby('Route')['Price'].mean()
average_journey_duration_by_route = df.groupby('Route')['Journey Duration'].mean()
average_delay_by_route = df.groupby('Route')['Delay'].mean()

print("Average Price by Route (Descending):")
display(average_price_by_route.sort_values(ascending=False))

print("\nAverage Journey Duration by Route (Descending):")
display(average_journey_duration_by_route.sort_values(ascending=False))

print("\nAverage Delay by Route (Descending):")
display(average_delay_by_route.sort_values(ascending=False))
```

➡ Average Price by Route (Descending):

	Price
Route	
Manchester Piccadilly to London Paddington	114.111111
Liverpool Lime Street to London St Pancras	104.774194
Liverpool Lime Street to London Euston	103.280766
Liverpool Lime Street to London Paddington	99.962963
Manchester Piccadilly to London St Pancras	99.562500
...	...
Liverpool Lime Street to Manchester Piccadilly	3.980680
Manchester Piccadilly to Liverpool Lime Street	3.740277
Manchester Piccadilly to Warrington	3.533333
London Euston to Oxford	2.562500
Birmingham New Street to Wolverhampton	1.875000

65 rows × 1 columns

**dtype:** float64

Average Journey Duration by Route (Descending):

	Journey Duration
Route	
Edinburgh Waverley to London Kings Cross	275.274510
London Kings Cross to Edinburgh Waverley	260.000000
Liverpool Lime Street to London Paddington	168.481481
Liverpool Lime Street to London St Pancras	150.000000
London Paddington to Liverpool Lime Street	150.000000
...	...
London Euston to Manchester Piccadilly	-120.561798
Reading to London Paddington	-130.675676
Liverpool Lime Street to Leeds	-150.000000
York to Durham	-779.740310
Birmingham New Street to Edinburgh	-1170.000000

65 rows × 1 columns

**dtype:** float64

Synthesize findings from the previous steps regarding popular routes and summarize the key insights about their characteristics.

```

# Synthesize findings from the previous steps regarding popular routes and summarize the key insights about their characteristics.

# Print the top 10 most popular routes by volume
top_10_routes = route_counts.head(10).index

# Print the top 10 most popular routes by volume
for route in top_10_routes:
    avg_price = average_price_by_route.get(route, 'N/A')
    avg_duration = average_journey_duration_by_route.get(route, 'N/A')
    avg_delay = average_delay_by_route.get(route, 'N/A')
    print(f"\nRoute: {route}")
    print(f" Average Price: {avg_price:.2f}" if isinstance(avg_price, (int, float)) else f" Average Price: {avg_price}")
    print(f" Average Journey Duration (minutes): {avg_duration:.2f}" if isinstance(avg_duration, (int, float)) else f" Average Journey Duration (minutes): {avg_duration}")
    print(f" Average Delay (minutes): {avg_delay:.2f}" if isinstance(avg_delay, (int, float)) else f" Average Delay (minutes): {avg_delay}")

# Print the top 10 most popular routes by volume
print("\nComparison of Popular Routes:")
print("While Manchester Piccadilly to Liverpool Lime Street is the most popular route by volume, other routes may have different characteristics.")
print("For example, routes to/from London stations often have higher average prices.")
print("Routes with longer distances tend to have longer journey durations and potentially higher average delays.")
print("The 'Journey Status' analysis showed that popular routes like Manchester Piccadilly to Liverpool Lime Street have a significant percentage of delayed journeys.")

# Print the overall takeaways
print("\nOverall Takeaways:")
print("- Route popularity is not directly correlated with average price, journey duration, or delay.")
print("- High-volume routes can still experience significant delays.")
print("- Factors like distance, station location (e.g., London), and ticket class likely play a larger role in determining price and duration than volume alone.")
```

➡ Summary of Key Insights about Popular Routes:

-----

The most popular route by number of purchases is: Manchester Piccadilly to Liverpool Lime Street with 4628 purchases.

Characteristics of Popular Routes (Top 10 by Purchase Count):

Route: Manchester Piccadilly to Liverpool Lime Street  
Average Price: 3.74  
Average Journey Duration (minutes): 35.14  
Average Delay (minutes): 5.14

Route: London Euston to Birmingham New Street  
Average Price: 11.96  
Average Journey Duration (minutes): 49.29  
Average Delay (minutes): 3.16

Route: London Kings Cross to York  
Average Price: 46.71  
Average Journey Duration (minutes): 40.41  
Average Delay (minutes): 0.54

Route: London Paddington to Reading  
Average Price: 16.88  
Average Journey Duration (minutes): 6.70  
Average Delay (minutes): 0.61

Route: London St Pancras to Birmingham New Street  
Average Price: 15.23  
Average Journey Duration (minutes): 65.89  
Average Delay (minutes): 0.00

Route: Liverpool Lime Street to Manchester Piccadilly  
Average Price: 3.98  
Average Journey Duration (minutes): -9.04  
Average Delay (minutes): 0.78

Route: Liverpool Lime Street to London Euston  
Average Price: 103.28  
Average Journey Duration (minutes): 138.69  
Average Delay (minutes): 26.01

Route: London Euston to Manchester Piccadilly  
Average Price: 85.68  
Average Journey Duration (minutes): -120.56  
Average Delay (minutes): 0.00

Route: Birmingham New Street to London St Pancras  
Average Price: 27.08  
Average Journey Duration (minutes): 80.00  
Average Delay (minutes): 0.00

Route: London Paddington to Oxford  
Average Price: 26.51  
Average Journey Duration (minutes): 90.00  
Average Delay (minutes): 0.00

Comparison of Popular Routes:  
While Manchester Piccadilly to Liverpool Lime Street is the most popular route by volume, other routes may have different character

## ✓ Analyze journey status distribution

Calculate and print the value counts and proportions of the 'Journey Status' column and create a bar plot to visualize the distribution.

```
journey_status_counts = df['Journey Status'].value_counts()
print("Number of journeys for each status:")
display(journey_status_counts)
journey_status_proportions = df['Journey Status'].value_counts(normalize=True)
print("\nProportion of journeys for each status:")
display(journey_status_proportions)
plt.figure(figsize=(8, 5))
sns.barplot(x=journey_status_proportions.index, y=journey_status_proportions.values)
plt.title('Distribution of Journey Status')
plt.xlabel('Journey Status')
plt.ylabel('Proportion')
plt.show()
```

➡ Number of journeys for each status:


count	
Journey Status	
On Time	27481
Delayed	2292
Cancelled	1880

dtype: int64

Proportion of journeys for each status:

proportion	
Journey Status	
On Time	0.868196
Delayed	0.072410
Cancelled	0.059394

dtype: float64



## ✓ Investigate reasons for delay

Calculate value counts and proportions for 'Reason for Delay' and visualize the proportions with a bar plot to understand the most common causes of delays.

```
reason_cols = [col for col in df.columns if col.startswith('Reason for Delay_')]
reason_for_delay_counts = df[reason_cols].sum().sort_values(ascending=False)
print("Number of occurrences for each reason for delay:")
display(reason_for_delay_counts)

reason_for_delay_proportions = reason_for_delay_counts / len(df)
print("\nProportion of each reason for delay:")
display(reason_for_delay_proportions)

plt.figure(figsize=(10, 6))
sns.barplot(x=reason_for_delay_proportions.index, y=reason_for_delay_proportions.values)
plt.title('Proportion of Journeys by Reason for Delay')
plt.xlabel('Reason for Delay')
plt.ylabel('Proportion')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

➡ Number of occurrences for each reason for delay:


	0
Reason for Delay_Unknown Reason	27481
Reason for Delay_Weather	1372
Reason for Delay_Signal Failure	970
Reason for Delay_Staff Shortage	809
Reason for Delay_Technical Issue	707
Reason for Delay_Traffic	314

dtype: int64

Proportion of each reason for delay:

	0
Reason for Delay_Unknown Reason	0.868196
Reason for Delay_Weather	0.043345
Reason for Delay_Signal Failure	0.030645
Reason for Delay_Staff Shortage	0.025558
Reason for Delay_Technical Issue	0.022336
Reason for Delay_Traffic	0.009920

dtype: float64



The 'Reason for Delay' column was one-hot encoded in a previous step and no longer exists. I need to use the one-hot encoded columns to analyze the reasons for delay. I will sum the one-hot encoded columns to get the counts for each reason, then calculate the proportions and plot them.

```
reason_cols = [col for col in df.columns if col.startswith('Reason for Delay_')]
reason_for_delay_counts = df[reason_cols].sum().sort_values(ascending=False)
print("Number of occurrences for each reason for delay:")
display(reason_for_delay_counts)
reason_for_delay_proportions = reason_for_delay_counts / len(df)
print("\nProportion of each reason for delay:")
display(reason_for_delay_proportions)
plt.figure(figsize=(10, 6))
sns.barplot(x=reason_for_delay_proportions.index, y=reason_for_delay_proportions.values)
plt.title('Proportion of Journeys by Reason for Delay')
plt.xlabel('Reason for Delay')
plt.ylabel('Proportion')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()
```

➡ Number of occurrences for each reason for delay:


	0
Reason for Delay_Unknown Reason	27481
Reason for Delay_Weather	1372
Reason for Delay_Signal Failure	970
Reason for Delay_Staff Shortage	809
Reason for Delay_Technical Issue	707
Reason for Delay_Traffic	314

dtype: int64

Proportion of each reason for delay:

	0
Reason for Delay_Unknown Reason	0.868196
Reason for Delay_Weather	0.043345
Reason for Delay_Signal Failure	0.030645
Reason for Delay_Staff Shortage	0.025558
Reason for Delay_Technical Issue	0.022336
Reason for Delay_Traffic	0.009920

dtype: float64



## ✓ Analyze delay by route

## ✓ Analyze delay reasons by station

```
# Calculate the proportion of each delay reason for each departure station
delay_reason_by_departure_station = df.groupby('Departure Station')[reason_cols].mean()
print("Proportion of each delay reason by Departure Station:")
display(delay_reason_by_departure_station)

# Calculate the proportion of each delay reason for each arrival station
delay_reason_by_arrival_station = df.groupby('Arrival Destination')[reason_cols].mean()
print("\nProportion of each delay reason by Arrival Station:")
display(delay_reason_by_arrival_station)

# Visualize the proportion of each delay reason by Departure Station for the top N stations
N = 10 # You can adjust this number
top_departure_stations = departure_counts.head(N).index
delay_reason_by_departure_station_top_N = delay_reason_by_departure_station.loc[top_departure_stations]
delay_reason_by_departure_station_top_N = delay_reason_by_departure_station_top_N.transpose()

plt.figure(figsize=(14, 8))
delay_reason_by_departure_station_top_N.plot(kind='bar', stacked=True, figsize=(14,8))
plt.title('Proportion of Delay Reasons by Top 10 Departure Stations')
plt.xlabel('Reason for Delay')
plt.ylabel('Proportion')
```

```
plt.xticks(rotation=45, ha='right')
plt.legend(title='Departure Station', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()

# Visualize the proportion of each delay reason by Arrival Station for the top N stations
top_arrival_stations = arrival_counts.head(N).index
delay_reason_by_arrival_station_top_N = delay_reason_by_arrival_station.loc[top_arrival_stations]
delay_reason_by_arrival_station_top_N = delay_reason_by_arrival_station_top_N.transpose()
```

```
plt.figure(figsize=(14, 8))
delay_reason_by_arrival_station_top_N.plot(kind='bar', stacked=True, figsize=(14,8))
plt.title('Proportion of Delay Reasons by Top 10 Arrival Stations')
plt.xlabel('Reason for Delay')
plt.ylabel('Proportion')
plt.xticks(rotation=45, ha='right')
plt.legend(title='Arrival Station', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.tight_layout()
plt.show()
```

```
Proportion of each delay reason by Departure Station:
/tmp/ipython-input-4183821481.py:2: FutureWarning: The default of observed=False is deprecated and will be changed to True in a futi
delay_reason_by_departure_station = df.groupby('Departure Station')[reason_cols].mean()
```

Departure Station	Reason for Delay_Signal Failure	Reason for Delay_Staff Shortage	Reason for Delay_Technical Issue	Reason for Delay_Traffic	Reason for Delay_Unknown Reason	Reason for Delay_Weather
Birmingham New Street	0.011236	0.026217	0.050562	0.001404	0.874532	0.036049
Bristol Temple Meads	0.000000	0.000000	0.000000	0.000000	1.000000	0.000000
Edinburgh Waverley	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000
Liverpool Lime Street	0.021487	0.026749	0.054155	0.016005	0.751370	0.130235
London Euston	0.044207	0.015543	0.008680	0.005854	0.896044	0.029673
London Kings Cross	0.034524	0.007567	0.016316	0.008040	0.919839	0.013715
London Paddington	0.016444	0.030889	0.018222	0.010222	0.906000	0.018222
London St Pancras	0.028527	0.012593	0.010280	0.007967	0.925469	0.015163
Manchester Piccadilly	0.044425	0.039646	0.018584	0.015398	0.825310	0.056637
Oxford	0.111111	0.000000	0.013889	0.006944	0.854167	0.013889
Reading	0.015152	0.013468	0.006734	0.003367	0.951178	0.010101
York	0.023732	0.055016	0.007551	0.008630	0.875944	0.029126

```
Proportion of each delay reason by Arrival Station:
/tmp/ipython-input-4183821481.py:7: FutureWarning: The default of observed=False is deprecated and will be changed to True in a futi
delay_reason_by_arrival_station = df.groupby('Arrival Destination')[reason_cols].mean()
```

Arrival Destination	Reason for Delay_Signal Failure	Reason for Delay_Staff Shortage	Reason for Delay_Technical Issue	Reason for Delay_Traffic	Reason for Delay_Unknown Reason	Reason for Delay_Weather
Birmingham New Street	0.038491	0.015242	0.008783	0.006458	0.905709	0.025316
Bristol Temple Meads	0.111111	0.000000	0.013889	0.006944	0.854167	0.013889
Cardiff Central	0.000000	0.000000	0.000000	0.000000	1.000000	0.000000
Coventry	0.000000	0.046154	0.000000	0.030769	0.923077	0.000000
Crewe	0.000000	0.000000	0.010363	0.000000	0.974093	0.015544

Calculate the proportion of delayed journeys for each route, sort the results, and print the top N routes with the highest proportion of delays.


```
delayed_journeys = df[df['Journey Status'] == 'Delayed']
delayed_proportion_by_route = delayed_journeys.groupby('Route').size() / df.groupby('Route').size()
delayed_proportion_by_route = delayed_proportion_by_route.sort_values(ascending=False)
N = 15
print(f"Top {N} Routes with Highest Proportion of Delayed Journeys:")
display(delayed_proportion_by_route.head(N))
```



Edinburgh Waverley	0.005616	0.011236	0.022472	0.000000	0.949438	0.011236
Leeds	0.250980	0.000000	0.000000	0.003922	0.745098	0.000000
Leicester	0.000067	0.000000	0.017804	0.005935	0.952522	0.020772
London Euston to York	1.000000					
Liverpool Lime Street	0.023453	0.007634	0.011151	0.013540	0.870569	0.033652
Edinburgh Waverley to London Kings Cross	1.000000					
York to Wakefield	1.000000	0	0.121889	0.038290	0.252712	0.469049
London Kings Cross	0.711030	0.607143	0.000000	0.000000	0.392857	0.000000
Manchester Piccadilly to London Euston	0.695652					
London Paddington	0.017004	0.045584	0.005698	0.005698	0.914530	0.011396
Liverpool Lime Street to London Paddington	0.450704					
Manchester Piccadilly to Leeds	0.000000	0.000000	0.008011	0.000000	0.927904	0.024032
Pancras	0.000000	0.000000				
Birmingham New Street to Manchester Piccadilly	0.428571					
Birmingham New Street to London Euston	0.352000	5	0.000000	0.000000	0.911765	0.000000
Manchester Piccadilly	0.0018145	0.000000	0.052671	0.009073	0.901210	0.010837
York to Doncaster	0.000000					
Oxford to Bristol Temple Meads	0.104167	5	0.000000	0.012658	0.835443	0.018987
Manchester Piccadilly to Nottingham	0.088608					
Nuneaton	0.000000	0.013699	0.004566	0.004566	0.936073	0.041096
Manchester Piccadilly to Liverpool Lime Street	0.076491	2	0.014446	0.016051	0.889246	0.035313
York to Durham	0.000000	0.062016	0.000000	0.028926	0.909091	0.004132
Peterborough	0.000000	0.057851				
London Euston to Birmingham New Street	0.057496	2	0.018878	0.009184	0.909439	0.015306
Sheffield	0.007353	0.000000	0.000000	0.000000	0.992647	0.000000
Stafford	0.000000	0.000000	0.000000	0.000000	0.978947	0.021053
Swindon	0.000000	0.017544	0.008772	0.004386	0.947368	0.021930

Create a bar plot to visualize the proportion of delayed journeys for the top N routes.

```
plt.figure(figsize=(12, 8))
sns.barplot(x=delayed_proportion_by_route.head(N).index, y=delayed_proportion_by_route.head(N).values)
plt.title(f'Top {N} Routes by Proportion of Delayed Journeys')
plt.xlabel('Route')
plt.ylabel('Proportion of Delayed Journeys')
plt.xticks(rotation=90, ha='right')
plt.tight_layout()
plt.show()
```



<Figure size 1400x800 with 0 Axes>

